

NinaivX — Dissertation Report

Ashwin Kumar Ramesh Kumar · MSc Artificial Intelligence, University of Sheffield

COMPLETED DRAFT — structure & completeness overview for supervisor review

8

CHAPTERS

3

APPENDICES

~15,240

WORDS

8

FIGURES

8

TABLES

29

REFERENCES

Status: the full report is **written and complete** (all chapters through Conclusion, plus appendices). It includes an empirical user study (n = 21) and a controlled fabrication experiment. Formatting follows the departmental guidelines (11pt Computer Modern, 1.15 line spacing, 37/25 mm binding margins, per-chapter figure/table numbering, centred chapter headers, numbered pages, glossary). Remaining before submission: four in-app screenshots, supervisor name and acknowledgements. This overview is generated from the LaTeX source; the written prose sits behind each heading below.

Contents

Chapter 1 Introduction

- 1.1 Problem and motivation
- 1.2 Aim
- 1.3 Objectives
- 1.4 Research questions
- 1.5 Contributions
- 1.6 Report structure

Chapter 2 Background and Literature Survey

- 2.1 Conversational AI: from rules to neural language models
- 2.2 Speech technology: synthesis, recognition and voice cloning
- 2.3 AI companions and the problem of loneliness
- 2.4 The digital afterlife: recreating the deceased
- 2.5 Hallucination and safety in generative systems
- 2.6 Ethics of the digital afterlife
- 2.7 Comparative summary of existing systems
- 2.8 Summary and identified gap

Chapter 3 Requirements and Analysis

- 3.1 From gaps to requirements
- 3.2 Must-have requirements
- 3.3 Should-have requirements
- 3.4 Could-have and won't-have requirements
- 3.5 Non-functional requirements
- 3.6 Ethical and safety requirements
- 3.7 Requirements summary and traceability

Chapter 4 Design

- 4.1 Design principles
- 4.2 Dual-mode architecture

- 4.3 The air-gap as a code-level guarantee
- 4.4 Temporal grounding of the legacy persona
- 4.5 Content and transparency policy
- 4.6 Memory design
- 4.7 Voice pipeline design
- 4.8 Data model
- 4.9 Security and privacy design
- 4.10 Reliability and error handling
- 4.11 Design decisions and alternatives considered

Chapter 5 Implementation

- 5.1 Backend overview
- 5.2 The agent
- 5.3 Prompts and safety guardrails
- 5.4 Memory summariser
- 5.5 The language-model layer
- 5.6 The voice pipeline
- 5.7 Authentication and sessions
- 5.8 Input validation and enforcement
- 5.9 The mobile application
- 5.10 Deployment and distribution
- 5.11 Development methodology and project management
- 5.12 Implementation challenges

Chapter 6 Testing

- 6.1 Testing strategy
- 6.2 Automated unit testing
- 6.3 Functional testing
- 6.4 Safety and robustness testing
 - 6.4.1 Worked examples
- 6.5 Limitations of the testing performed

Chapter 7 Evaluation

- 7.1 Evaluation approach
- 7.2 User-study method
- 7.3 Participants
- 7.4 Quantitative results
- 7.5 Qualitative results
- 7.6 Controlled fabrication evaluation
 - 7.6.1 Method
 - 7.6.2 Results
- 7.7 Discussion: answering the research questions
- 7.8 Positioning against existing systems
- 7.9 Ethical reflection
- 7.10 Legal, social and professional considerations
- 7.11 Threats to validity and limitations

Chapter 8 Conclusion

- 8.1 Summary
- 8.2 Contributions and findings
- 8.3 Reflection against the aims and research questions
- 8.4 Future work
- 8.5 Reflection and lessons learned
- 8.6 Closing remarks

Appendix A Glossary of abbreviations

Appendix B Ethics approval and study protocol

Appendix C User-study questionnaire

List of Figures (8)

- **Figure 4.1** — The dual-mode architecture of NinaivX: the mobile client speaks only to the API layer, which authenticates the request, confirms persona own
- **Figure 4.2** — The agent graph: a deterministic supervisor routes a deceased persona to the air-gapped node (a single edge to end, no tools bound) and a co
- **Figure 4.3** — The NinaivX data model: a User owns many Personas, and each Persona has exactly one Conversation (which holds the rolling summary).
- **Figure 5.1** — Key screens of the NinaivX mobile application: (a) sign-up and log-in, (b) the three-step create-persona wizard, (c) the text-chat interface
- **Figure 5.2** — The end-to-end workflow of NinaivX, from installing the app through signing up, creating a persona, and holding a text or voice conversation
- **Figure 5.3** — Project timeline: the major phases of NinaivX from literature survey and requirements through design, implementation, testing, evaluation an
- **Figure 7.1** — Distribution of user-study responses across the twelve questionnaire items (n=21); segment labels give the number of respondents at each lev
- **Figure 7.2** — Fabrication rate by question type: the constrained NinaivX design versus an unconstrained baseline on the same model. Temporal grounding all

List of Tables (8)

- **Table 2.1** — How NinaivX compares with representative existing systems (yes, partial, × no).
- **Table 3.1** — Gaps in existing work, derived from the literature survey.
- **Table 3.2** — Must-have requirements (= met). All fifteen are met.
- **Table 3.3** — Should-have requirements.
- **Table 3.4** — Non-functional requirements and how they are addressed.
- **Table 6.1** — Summary of the automated test suite.
- **Table 7.1** — User-study results by item (n=21; five-point scale).
- **Table 7.2** — Fabrication rates by question type (three samples per probe; same model, grounding varied).